

4. The Abstraction: The Process

Operating System: Three Easy Pieces

How to provide the illusion of many CPUs?

▣ CPU virtualizing

- ◆ The OS can promote the illusion that many virtual CPUs exist.
- ◆ **Time sharing**: Running one process, then stopping it and running another
 - The potential cost is **performance**.

A process is a running program.

- Comprising of a process:
 - ◆ Memory (address space)
 - Instructions
 - Data section
 - ◆ Registers (processor architectural state (?))
 - Program counter
 - Stack pointer
 - ...

Process API

- ▣ These APIs are available on any modern OS.
 - ◆ **Create**
 - Create a new process to run a program
 - ◆ **Destroy**
 - Halt a runaway process
 - ◆ **Wait**
 - Wait for a process to stop running
 - ◆ **Miscellaneous Control**
 - Some kind of method to suspend a process and then resume it
 - ◆ **Status**
 - Get some status info about a process
 - ◆ ...

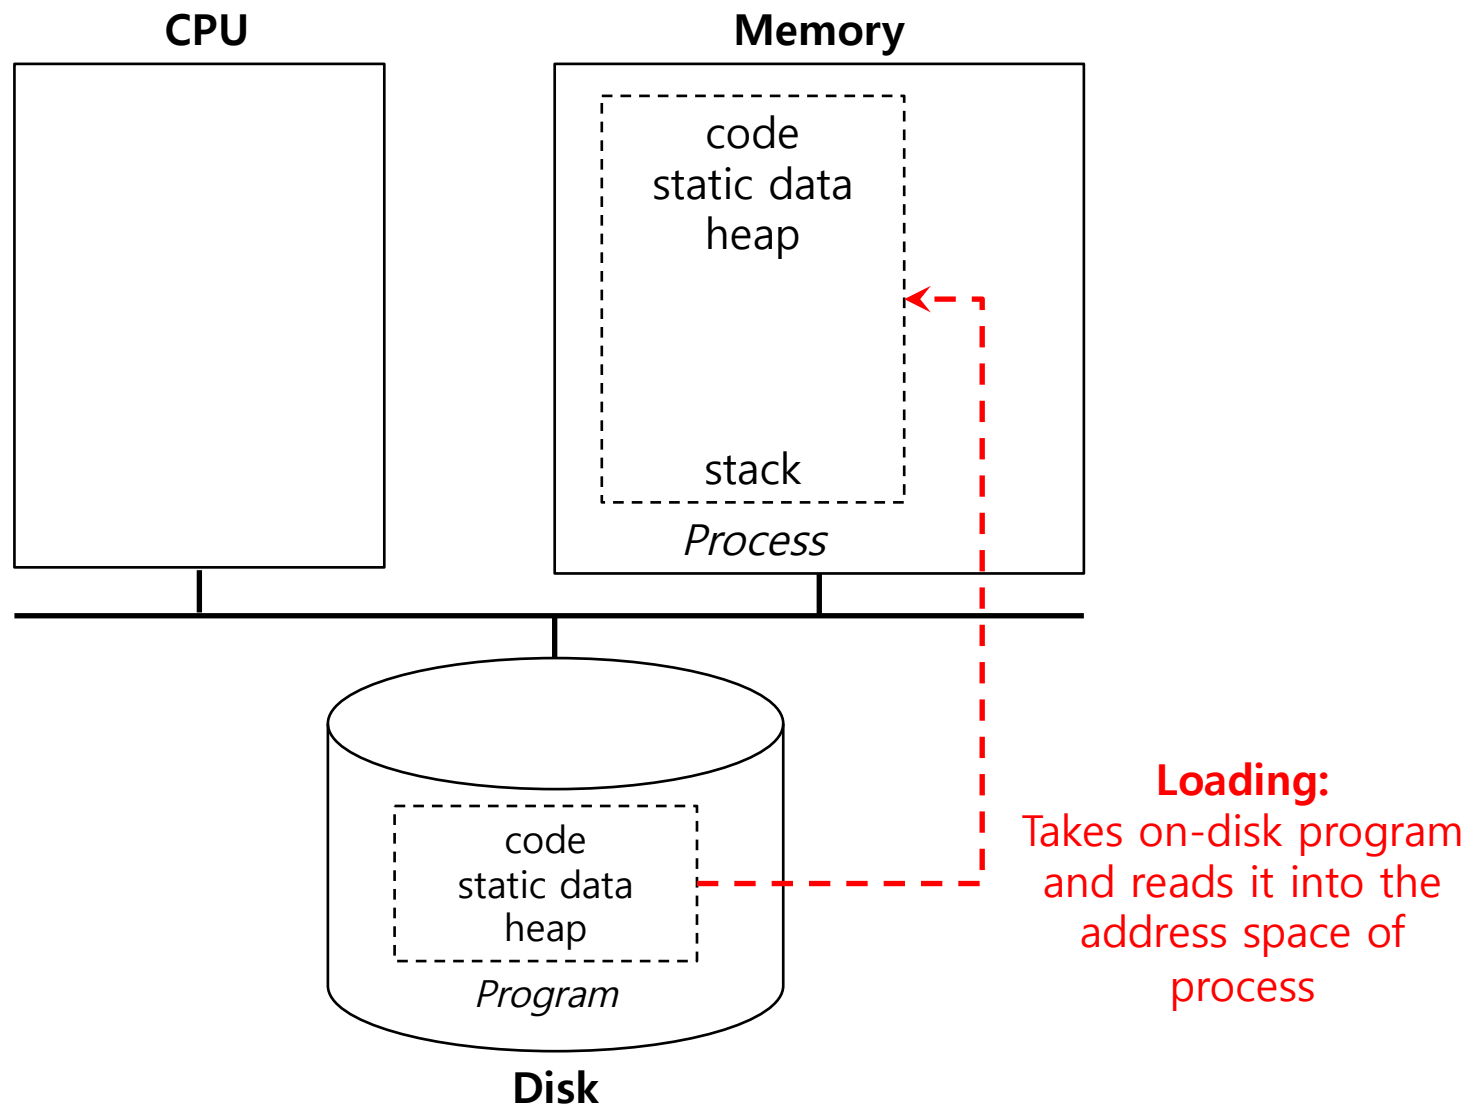
Process Creation

1. **Load** a program code into memory, into the address space of the process.
 - ◆ Programs initially reside on disk in *executable format (code + static data)*.
 - ◆ OS perform the loading process **lazily**.
 - Loading pieces of code or data only as they are needed during program execution.
2. The program's run-time **stack** is allocated.
 - ◆ Use the stack for *local variables, function parameters, and return address*.
 - ◆ Initialize the stack with arguments → `argc` and the `argv` array of `main()` function

Process Creation (Cont.)

3. The program's **heap** is created.
 - ◆ Used for explicitly requested dynamically allocated data.
 - ◆ Program request such space by calling `malloc()` and free it by calling `free()`.
4. The OS do some other initialization tasks.
 - ◆ input/output (I/O) setup
 - Each process by default has three open file descriptors.
 - Standard input, output and error
5. **Start the program** running at the entry point, namely `main()`.
 - ◆ The OS *transfers control* of the CPU to the newly-created process.

Loading: From Program To Process



Process States

- A process can be one of three states.

- ◆ **Running**

- A process is running on a processor.

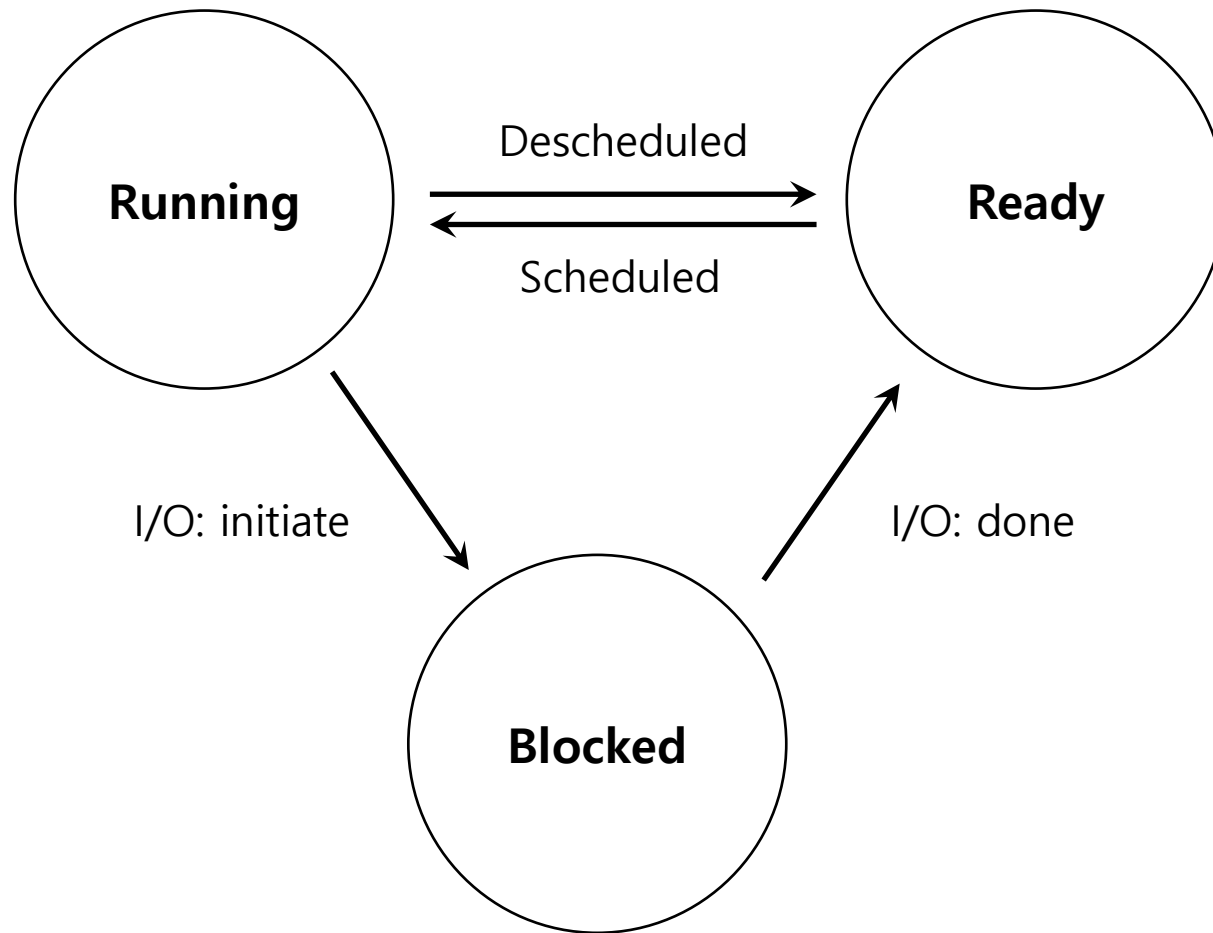
- ◆ **Ready**

- A process is ready to run but for some reason the OS has chosen not to run it at this given moment.

- ◆ **Blocked**

- A process has performed some kind of operation.
 - When a process initiates an I/O request to a disk, it becomes blocked and thus some other process can use the processor.

Process State Transition



Data structures

- The OS has **some key data structures** that track various relevant pieces of information.
 - ◆ **Process list**
 - Ready processes
 - Blocked processes
 - Current running process
 - ◆ **Register context**
- PCB(Process Control Block)
 - ◆ A C-structure that contains information **about each process**.

Example) The xv6 kernel Proc Structure

```
// the registers xv6 will save and restore
// to stop and subsequently restart a process
struct context {
    int eip;    // Index pointer register
    int esp;    // Stack pointer register
    int ebx;    // Called the base register
    int ecx;    // Called the counter register
    int edx;    // Called the data register
    int esi;    // Source index register
    int edi;    // Destination index register
    int ebp;    // Stack base pointer register
};

// the different states a process can be in
enum proc_state { UNUSED, EMBRYO, SLEEPING,
                  RUNNABLE, RUNNING, ZOMBIE };
```

Example) The xv6 kernel Proc Structure (Cont.)

```
// the information xv6 tracks about each process
// including its register context and state
struct proc {
    char *mem;           // Start of process memory
    uint sz;             // Size of process memory
    char *kstack;        // Bottom of kernel stack
                        // for this process

    enum proc_state state; // Process state
    int pid;              // Process ID
    struct proc *parent;  // Parent process
    void *chan;           // If non-zero, sleeping on chan
    int killed;           // If non-zero, have been killed
    struct file *ofile[NOFILE]; // Open files
    struct inode *cwd;    // Current directory
    struct context context; // Switch here to run process
    struct trapframe *tf; // Trap frame for the
                        // current interrupt
};
```

- Disclaimer: This lecture slide set is used in AOS course in University of Cantabria. Was initially developed for Operating System course in Computer Science Dept. at Hanyang University. This lecture slide set is for OSTEP book written by Remzi and Andrea Arpaci-Dusseau (at University of Wisconsin)